

Predicting California House Prices

May 13, 2024

```
[2]: import tensorflow as tf
import pandas as pd
```

0.1 The Data

Import the cal_housing_clean.csv file with pandas. Separate it into a training (70%) and testing set(30%).

```
[3]: df=pd.read_csv('/Users/sandinatatu/Desktop/Tensorflow-Bootcamp-master/
↳02-TensorFlow-Basics/cal_housing_clean.csv')

df.head()
```

```
[3]:   housingMedianAge  totalRooms  totalBedrooms  population  households  \
0                41.0        880.0          129.0        322.0         126.0
1                21.0       7099.0         1106.0       2401.0        1138.0
2                52.0       1467.0          190.0        496.0         177.0
3                52.0       1274.0          235.0        558.0         219.0
4                52.0       1627.0          280.0        565.0         259.0

      medianIncome  medianHouseValue
0          8.3252      452600.0
1          8.3014      358500.0
2          7.2574      352100.0
3          5.6431      341300.0
4          3.8462      342200.0
```

Separate your features and target data (medianHouseValue) into training and testing sets, with 33% reserved for testing.

```
[4]: X = df.iloc[:, :6]
y=df['medianHouseValue']

import numpy as np
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
↳random_state=42)
```

0.2 Scale the Feature Data

Use sklearn preprocessing to create a MinMaxScaler for the feature data. Fit this scaler only to the training data. Then use it to transform X_test and X_train.

```
[5]: from sklearn.preprocessing import MinMaxScaler
      scaler = MinMaxScaler().fit(X_train)
      X_train = scaler.transform(X_train)
      X_test = scaler.transform(X_test)
```

0.3 Fit a Densely Connected Neural Network to the Training Data

Construct a Densely Connected Neural Network with 3 hidden layers, each having 6 neurons to predict the Median House Value.

```
[7]: model = tf.keras.models.Sequential()
      model.add(tf.keras.Input(shape=(6,)))
      model.add(tf.keras.layers.Dense(6, activation='relu'))
      model.add(tf.keras.layers.Dense(6, activation='relu'))
      model.add(tf.keras.layers.Dense(6, activation='relu'))
      model.add(tf.keras.layers.Dense(1, activation='linear'))

      learning_rate = 0.01
      optimizer = tf.keras.optimizers.SGD(learning_rate=learning_rate)

      model.compile(loss='mse',
                    optimizer=optimizer)

      model.fit(X_train, y_train, epochs=100, batch_size=10, shuffle=True, verbose=0)
```

```
[7]: <keras.src.callbacks.history.History at 0x109cf9650>
```

0.4 Compute the RMSE on the Test Data

```
[8]: mse = model.evaluate(X_test, y_test, verbose=0)

      print("Root Mean Squared Error on Test Data:", mse**0.5)
```

Root Mean Squared Error on Test Data: 115438.0119024925

```
[ ]:
```